AssignmentB

Martijn Koster, William Schaafsma, Martijn van Dam, Victor Hovius

3/21/2022

Contents

1   Introduction

The matter of interest for this assignment will be the impact that incomplete data (**observed data**) have on our inferences compared to the inferences we make with complete data (**true data**). To investigate the effect that missing values have on model inferences, we will build a multiple regression model.

Firstly, we provide descriptive statistics and correlations. In table 1 we compare the head of the observed data and the true data. Additionally, in Table 3 the means and variances are compared. With regard to correlations, we present two correlations matrices; one for the observed data (Table 6) and the other for the true data Table (7).

Secondly, we present our multiple regression model in table 8. Our model consists of the outcome variable: *active heart rate* and the predictors: *age* and *smoke*. We also included an interaction effect between *bmi* and *sex*. The first three columns reflect the observed data, whereas the latter reflect the true data.

The research question we try to answer in accordance with our model is: *What influence does a person's age, smoking habits, sex and bmi have on their heart rate during exercise?*

Thirdly, we start by inspecting the missing values. Then, we try to find out where the missing values occur. In Figure 1 we begin by giving a global overview of the missingness.

Afterwards, we perform t-tests on the variables containing missing values to check the type of missingness, either MNAR, MAR, or MCAR. We also provide plots here to visualize where the missing values occur.

Lastly, we use multiple imputation to solve the missingness. How we are doing it, and why will be further elaborated in Section 5

2   Observed vs True data

In this section we will compare the observed with the true data set.

Table 1: First Five Cases of Observed Data

age	smoke	sex	intensity	active	rest	height	weight	bmi
42	no	female	high	NA	75	NA	NA	22.4
31	NA	male	low	NA	62	NA	NA	23.8
36	no	male	low	109	76	182	78.0	23.5
31	no	female	low	78	62	164	53.9	20.0
42	no	male	low	NA	66	189	NA	23.4

Table 2: First Five Cases of True Data

age	smoke	sex	intensity	active	rest	height	weight	bmi
42	no	female	high	94	75	161	58.1	22.4
31	no	male	low	86	62	184	80.6	23.8
36	no	male	low	109	76	182	78.0	23.5
31	no	female	low	78	62	164	53.9	20.0
42	no	male	low	103	66	189	83.6	23.4

2.1   Descriptives

The means and the variance of the variables age and rest remain unchanged, as they have no missing values.

Table 3: Means and variances in true and observed dataset

Variables	M obs	M true	var obs	var true	N obs	N true
Age	38.52	38.52	149.73	149.73	306	306
Active	92.58	93.13	383.05	378.04	183	306
Rest	69.83	69.83	120.78	120.78	306	306
Height	174.50	173.99	100.66	105.29	214	306
Weight	73.28	73.58	270.28	274.85	132	306
Bmi	24.11	24.06	12.91	13.38	213	306

Note.

obs = Observed Dataset, true = True Dataset

The means of the variables active, weight, and bmi are somewhat higher for the observed data set compared to the true data set. For height, the mean of the observed data set is somewhat smaller compared to the true data set.

Furthermore, the variance of the variable active in the observed data set is .01 lower than the true data set, and thus almost entirely unaffected. However the variables height, weight, and bmi have greater variance in the true data set than the observed data set. This implies that the missingness causes an underestimation of the variance.

2.2 Categorical variables

The categorical variables in both data sets are *smoke*, *sex* and *intensity*. *Smoke* and *sex* both have two levels ("no" and "yes" for smoke and "male" and "female" for sex), while *intensity* has three levels ("low", "moderate", and "high"). Despite differences in the number of observed values between the data sets, differences between groups remain unchanged. For example, there are more males than females in both data sets and more non-smokers than smokers. Also, in both data sets, more males reported smoking than females. The most frequently reported workout intensity for both males and females in the two data sets is moderate, followed by low and high.

Table 4: proportion table of categorical variables in observed data

sex	smoke	intensity	Freq
male	no	high	0.040
female	no	high	0.060
male	yes	high	0.065
female	yes	high	0.056
male	no	moderate	0.113
female	no	moderate	0.149
male	yes	moderate	0.085
female	yes	moderate	0.073
male	no	low	0.165
female	no	low	0.129
male	yes	low	0.048
female	yes	low	0.016

Table 5: proportion table of categorical variables in true data

sex	smoke	intensity	Freq
male	no	high	0.046
female	no	high	0.059
male	yes	high	0.075
female	yes	high	0.046
male	no	moderate	0.108
female	no	moderate	0.147
male	yes	moderate	0.085
female	yes	moderate	0.062
male	no	low	0.190
female	no	low	0.124
male	yes	low	0.046
female	yes	low	0.013

2.3 Correlations

As shown in Table 6 and Table 7, the correlations between the variables of the observed data set are slightly different from the correlations between variables of the true data. Although most of the correlations are almost identical, a few correlations are negative in the observed data and positive in the true data. This effect also occurs vice versa. In example, the correlation between the variables smoke and age of the observed data set is positive ($r = 0.01$), albeit almost 0. In contrast, the correlation for these variables in the true data set is negative ($r = $ -0.05). However, the impact of missing data on the correlations appears to be minor, as the difference in correlation coefficients between the two data sets is negligible. Although some correlations differ in valency between the data sets, the correlation coefficients remain close to 0 and thus, do not distort inferences made with the observed data set.

Table 6: Correlations of observed data

	age	smoke	sex	intensity	active	rest	height	weight	bmi
age	1.00	0.01	-0.17	0.21	-0.49	-0.39	0.19	0.25	0.18
smoke	0.01	1.00	-0.09	-0.29	0.15	0.23	0.18	0.18	0.18
sex	-0.17	-0.09	1.00	-0.09	0.11	0.06	-0.73	-0.68	-0.42
intensity	0.21	-0.29	-0.09	1.00	-0.37	-0.55	0.13	0.12	0.02
active	-0.49	0.15	0.11	-0.37	1.00	0.56	0.00	0.01	0.05
rest	-0.39	0.23	0.06	-0.55	0.56	1.00	-0.20	-0.12	0.06
height	0.19	0.18	-0.73	0.13	0.00	-0.20	1.00	0.78	0.34
weight	0.25	0.18	-0.68	0.12	0.01	-0.12	0.78	1.00	0.88
bmi	0.18	0.18	-0.42	0.02	0.05	0.06	0.34	0.88	1.00

Table 7: Correlations of true data

	age	smoke	sex	intensity	active	rest	height	weight	bmi
age	1.00	-0.05	-0.17	0.21	-0.54	-0.39	0.20	0.23	0.20
smoke	-0.05	1.00	-0.11	-0.31	0.18	0.27	0.17	0.25	0.24
sex	-0.17	-0.11	1.00	-0.09	0.09	0.06	-0.72	-0.69	-0.47
intensity	0.21	-0.31	-0.09	1.00	-0.37	-0.55	0.12	0.06	0.01
active	-0.54	0.18	0.09	-0.37	1.00	0.61	-0.10	0.02	0.09
rest	-0.39	0.27	0.06	-0.55	0.61	1.00	-0.15	-0.04	0.05
height	0.20	0.17	-0.72	0.12	-0.10	-0.15	1.00	0.77	0.36
weight	0.23	0.25	-0.69	0.06	0.02	-0.04	0.77	1.00	0.87
bmi	0.20	0.24	-0.47	0.01	0.09	0.05	0.36	0.87	1.00

3 Regression

3.1 Answering the research question

When examining Table 8 *Regression analysis of True and Observed data* we observe several differences in the beta coefficients, standard error, and p-values. The table contains variables with missing values and an interaction effect. Although almost all beta coefficients are nearly equal, the beta coefficients of the observed data set are systematically underestimated. This underestimation is especially the case for *sexfemale*, as the difference between the beta coefficients is almost 9.0. Making inferences based on the observed data set would lead to underestimating the effect of sex on active hear rate. Regarding the standard errors, missing data caused these parameters of the observed data set to be systematically overestimated. Larger standard errors contribute to the possibility of making a type II error, as is the case in our data set. For example, the larger standard errors in the observed data set might have played a role in the variables *sexfemale* and the interaction *bmi:sexfemale* turning non-significant. These variables would wrongly be neglected when making inferences with the model based on the observed data. Concluding, the missing data causes the standard errors to be greater, resulting in less accurate beta coefficients. Moreover, some p-values turn out non-significant, caused by underestimated beta coefficients. The model based on observed data leads thus to inaccurate inferences.

Table 8: Regression analysis of True (N=306) and Observed Data (N=155)

	Observed Data			True Data		
	b	SE	p	b	SE	p
(Intercept)	78.444	14.34	0.000	80.384	9.03	0.000
age	-0.809	0.11	0.000	-0.883	0.07	0.000
bmi	1.681	0.55	0.003	1.776	0.35	0.000
sexfemale	32.756	20.78	0.117	43.460	14.16	0.002
smokeyes	1.615	2.91	0.580	3.516	1.99	0.078
bmi:sexfemale	-1.131	0.88	0.199	-1.674	0.60	0.006

4 Missingness

There are 540 missing values. 0 for age, 0 for sex, 0 for intensity, 0 for rest, 58 for smoke, 92 for height, 93 for bmi, 123 for active, and 174 for weight. Moreover, there are 132 completely observed rows, 15 rows with one missing value, 37 rows with two missing values, 52 rows with three missing values, 55 rows with four missing values, 15 rows with five missing values. The missingness in the data is non-monotone because the variable with the least missing values (*smoke*) has observed values for other variables with more missingness (e.g., *smoke* and *bmi*). The missingness would be monotone if the variable with the least missing values (*smoke*), would have missing values on all other variables with more missingness (e.g., *height*). Interestingly, a monotone pattern is only the case for *smoke* and *weight*.

Figure 1: pattern of the missingness

4.1 Looking for the missingness

In this section, we will investigate whether the mean of the missing values differs significantly from the mean of the observed values. This will be done do by using a paired sampled t-test for the numeric variables. To compare the mean of the missing values with the true values, we computed a logical vector for each vector that has missing observations. The missingness vectors have the value `TRUE` for all missing entries and `FALSE` for all observed entries. These missingness vectors will be used as a grouping variable in the true data set to compare the missing values with the observed values. For smoke, which is a categorical variable, we will use a x^2 test. For all variables, the missing values have a similar distribution as the observed values. However, the distribution for the variable smoke is not shown, as this is a categorical variable and does thus not have a distribution. The means of the variables from both data sets are marginally different, but the differences are non-significant, neither for smoke. Hence, the missing values are similar to the observed values.

4.2 Missingness of weight

Looking at the differences of the missingness of weight in Table 9, no significant differences can be found in the means of the numeric variables in the data. However, the difference in mean of active is relatively high. It might be that this difference is not significant due to the low amount of observations of active when weight is missing ($N = 36$). |Considering the categorical data, the missingness of weight on sex has no significant difference, where $x^2 = 0$, $p = 1$. For the missingness of weight on smoke no significant difference was found also $x^2 = 0.036$, $p = 0.848$. Lastly, the missingness of weight on intensity is not significant: $x^2 = 2.589$, $p = 0.274$.
| All results are non-significant; hence weight is not missing at random.

Table 9: Difference in means of missingness of Weight

variables	$M obs	M true	t	p
rest	69.56	70.17	-0.482	0.630
age	38.16	38.98	-0.590	0.556
height	174.28	175.39	-0.639	0.525
bmi	24.11	24.12	-0.012	0.990
active	91.54	96.81	-1.440	0.156

The three bar plots in Figure 2 show a visualization of how the missing data in the categorical columns is divided. The first plot shows us that there is almost no difference between missing values in weight for being a man or female in the sex column. The second plot also shows that there is almost no difference between missing values in the weight column for smokers and non-smokers in the smoke column. The third column shows that how lower the intensity is the less missing values in weight you can expect.

(a) Sex   (b) Smoke   (c) Intensity
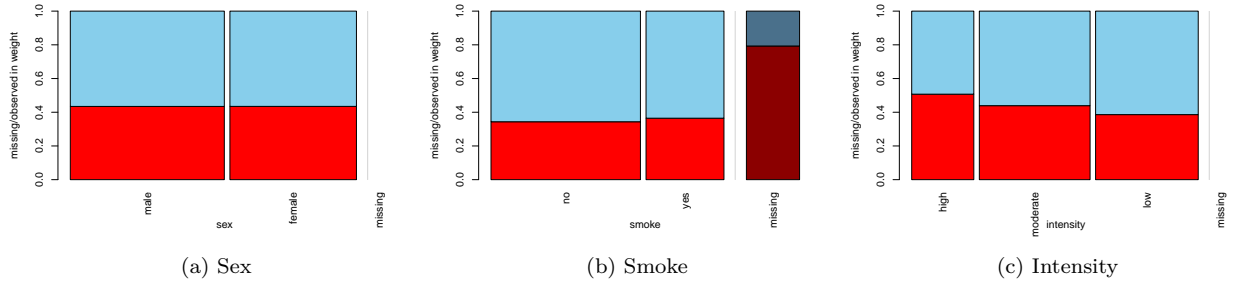
Figure 2: Looking whether the missingness of weight is MAR

The five scatterplots in Figure 3 show a visualization of how the missing data is divided in the rest of the columns. In the first two plots between weight and rest or age is a clear trend where all the values with a low weight are missing, and everything above that is not. The two plots after that between weight and height or bmi show the same thing, but also a cluster of missing values when both columns have low values. The last column between weight and active shows a clear trend where low values for either column results in missing values with a cluster where both columns have low values.
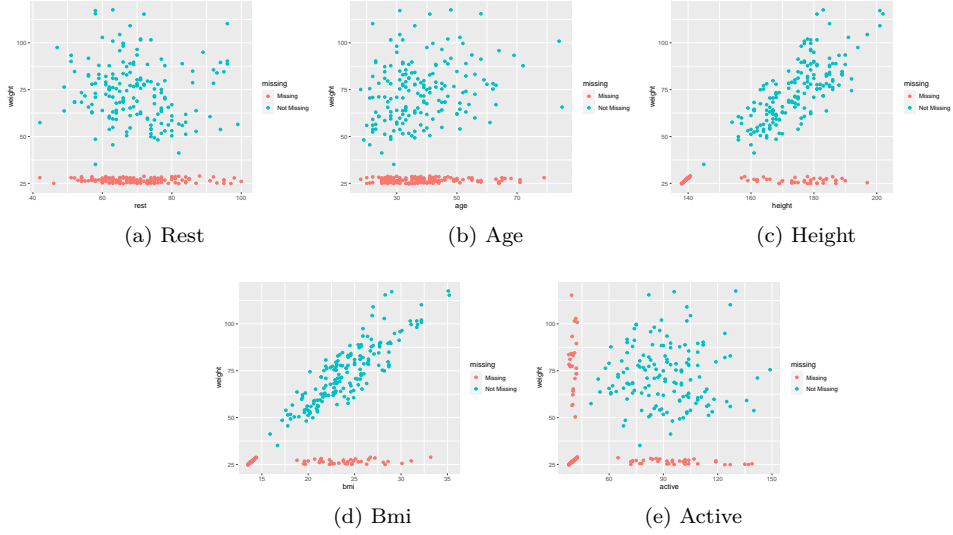
Figure 3: Looking whether the missingness of weight is MAR

4.3 Missingness of height

Looking at the differences of the missingness of height in Table 10, no significant differences can be found in the means of the numeric variables in the data. Similar to the missingness of weight, the difference in mean of `active` is relatively high. Again, it might be that this difference is not significant due to the low amount of observations of `active` when height is missing ($N = 40$). Considering the categorical data, the missingness of height on sex has no significant difference, where $x^2 = 0$, $p = 1$. For the missingness of height on smoke no significant difference was found also $x^2 = 0.111$, $p = 0.739$. Lastly, the missingness of weight on intensity is not significant: $x^2 = 3.563$, $p = 0.168$.
All results are insignificant, hence height is not missing at random.

Table 10: Difference in means of missingness of Height

variables	M obs	M true	t	p
rest	69.93	69.60	0.242	0.809
age	38.66	38.18	0.320	0.749
bmi	24.11	24.11	-0.012	0.990
active	91.74	99.05	-1.535	0.137

The three bar plots in Figure 4 show a visualization of how the missing data in the categorical columns is divided. The first plot shows us almost no difference between missing values in height for being a man or female in the sex column. The second plot also shows virtually no difference between missing values in the height column for smokers and non-smokers in the smoke column. The third column shows almost no difference between missing values in the height column for high and moderate-intensity but less missing values in the low-intensity category.

(a) Sex   (b) Smoke   (c) Intensity
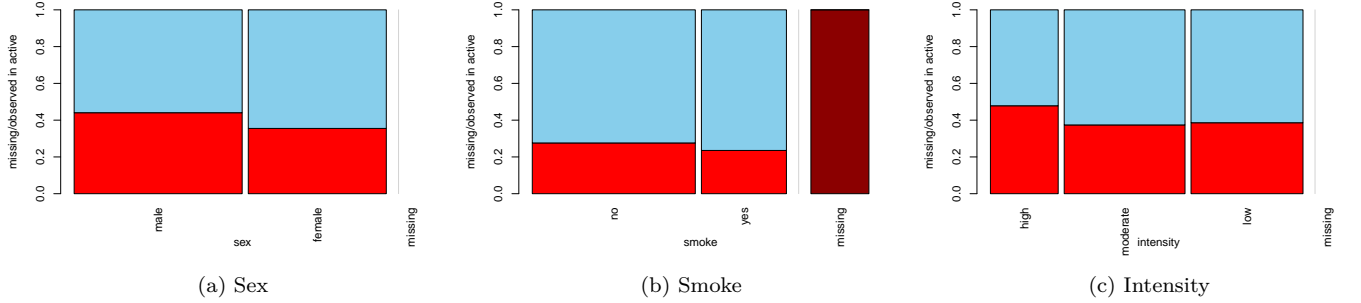
Figure 4: Looking whether the missingness of height is MAR

The four scatterplots in Figure 5 show how the missing data is divided into the rest of the columns. There is a clear trend in the first two plots between height and rest or age where all the values with a low height are missing and everything above that is not. The two plots after that between height and bmi or active show a clear trend where low values for either column result in missing values with a cluster where both columns have low values.

(a) Rest   (b) Age   (c) Bmi
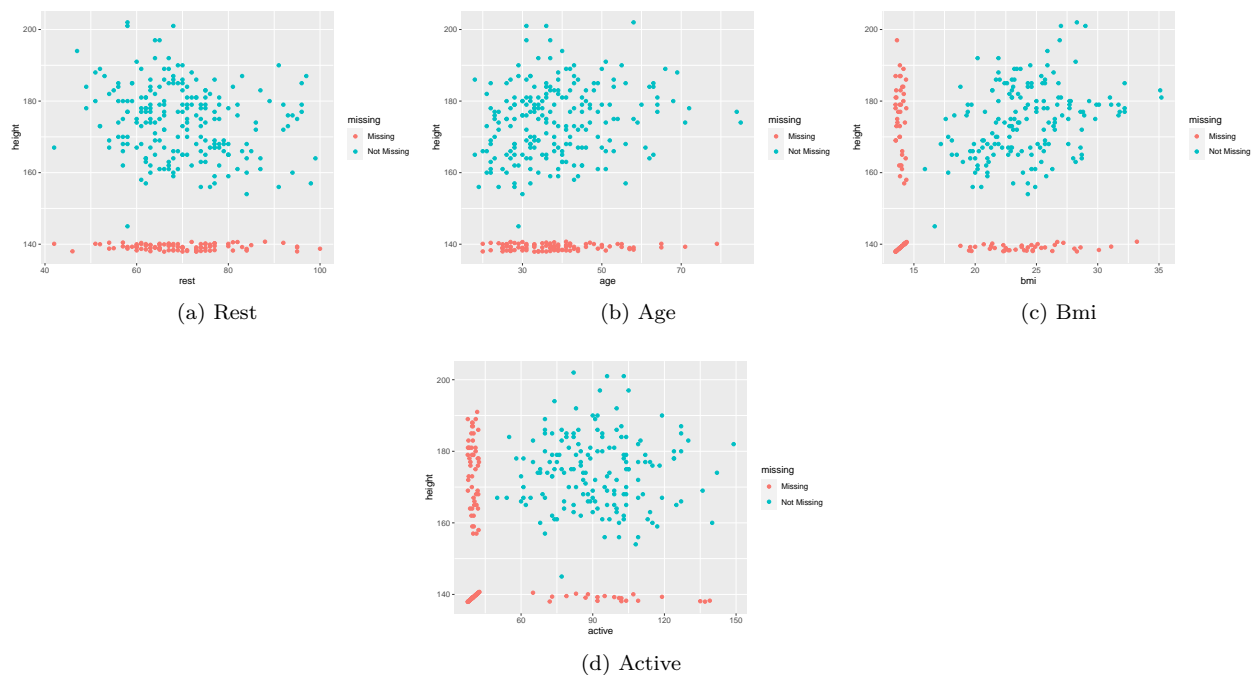
(d) Active

Figure 5: Looking whether the missingness of height is MAR

4.4 Missingness of Active

Looking at the differences of the missingness of active in Table 11, no significant differences can be found in the means of the numeric variables in the data. However, both `bmi` ($p = 0.062$) and `weight` ($p = 0.059$) are relatively close to the significance threshold of 0.05. Considering the categorical data, the missingness of active on sex has no significant difference, where $x^2 = 1.957$, $p = 0.162$. For the missingness of active on smoke no significant difference was found also $x^2 = 0.293$, $p = 0.589$. Lastly, the missingness of weight on intensity is not significant: $x^2 = 2.193$, $p = 0.334$. All results are insignificant, hence active is not considered to be missing at random.

Table 11: Difference in means of missingness of Active

variables	M obs	M true	t	p
rest	69.02	71.03	-1.558	0.120
age	37.96	39.35	-0.963	0.337
height	174.41	174.77	-0.232	0.817
bmi	23.83	24.85	-1.883	0.062
weight	72.94	79.38	-1.948	0.059

The three bar plots in Figure 6 show a visualization of how the missing data in the categorical columns is divided. The first plot shows us that the female category in sex has less missing values in the active column than the male category. The second column shows that smokers have fewer missing values than non-smokers in the active column. The third column shows almost no difference between missing values in the active column for the moderate and low-intensity category but more missing values in the high-intensity category.

(a) Sex   (b) Smoke   (c) Intensity

Figure 6: Looking whether the missingness of active is MAR

The five scatterplots in Figure 7 show a visualization of how the missing data is divided into the rest of the columns. In the first two plots between active and rest or age, there is a clear trend where all the values with a low active are missing, and everything above that is not. The three plots after that between active and height, bmi, or weight show a clear trend where low values for either column result in missing values with a cluster where both columns have low values.
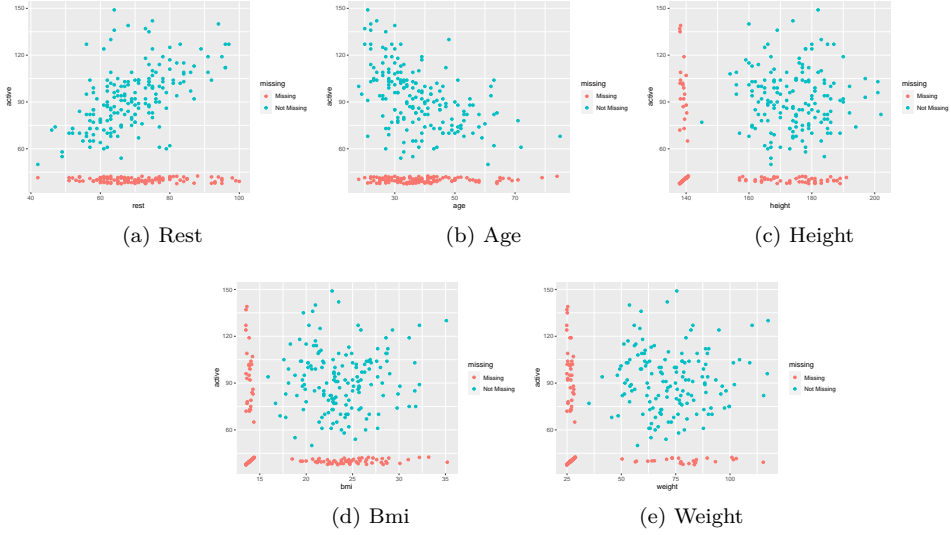
(a) Rest (b) Age (c) Height

(d) Bmi (e) Weight

Figure 7: Looking whether the missingness of active is MAR

4.5 Missingness of Bmi

Looking at the differences of the missingness of bmi in Table 12, no significant differences can be found in the means of the numeric variables in the data.

Considering the categorical data, the missingness of bmi on sex has no significant difference, where $x^2 = 0.019$, $p = 0.889$. For the missingness of bmi on smoke no significant difference was found also $x^2 = 0$, $p = 1$. Lastly, the missingness of bmi on intensity is not significant: $x^2 = 1.476$, $p = 0.478$.

All results are insignificant, hence bmi is considered not to be missing at random.

Table 12: Difference in means of missingness of Bmi

variables	M obs	M true	t	p
rest	69.84	69.81	0.021	0.983
age	38.35	38.90	-0.368	0.713
height	174.28	175.39	-0.639	0.525
active	92.12	95.11	-0.717	0.478

The three bar plots in Figure 8 show a visualization of how the missing data in the categorical columns is divided. The first plot shows us almost no difference between missing values in bmi for being a man or female in the sex column. The second plot also indicates practically no difference between missing values in the bmi column for smokers and non-smokers in the smoke column. The third column shows almost no difference between missing values in the bmi column for the moderate and low-intensity category, but more missing values in the high-intensity category.

(a) Sex   (b) Smoke   (c) Intensity

Figure 8: Looking whether the missingness of bmi is MAR

The four scatterplots in Figure 9 show a visualization of how the missing data is divided into the rest of the columns. In the first two plots between bmi and rest or age there is a clear trend where all the values with a low bmi are missing and everything above that is not. The two plots after that between bmi and height or active show a clear trend where low values for either column result in missing values with a cluster where both columns have low values.

(a) Rest   (b) Age   (c) Height

(d) Active
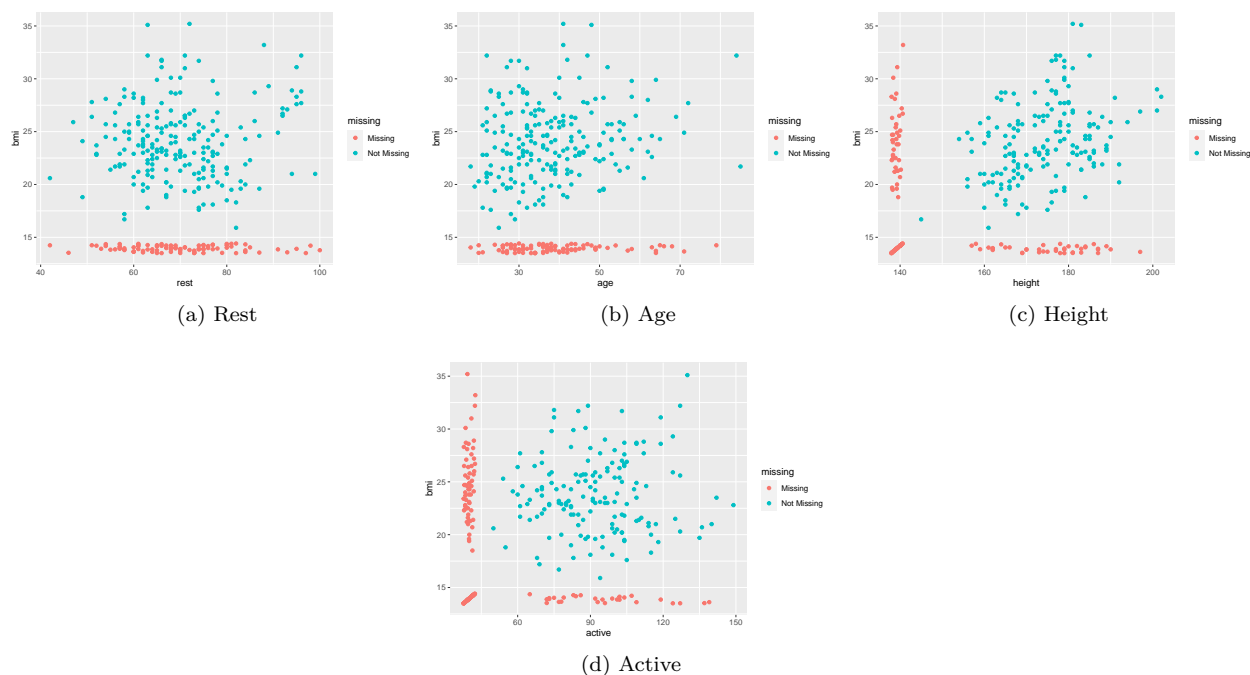
Figure 9: Looking whether the missingness of bmi is MAR

4.6 Missingness of Smoke

Looking at the differences of the missingness of smoke in Table 13, no significant differences can be found in the means of the numeric variables in the data.

Considering the categorical data, the missingness of weight on sex has a significant difference, where $x^2 = 5.037$, $p = 0.025$. The missingness of smoke on intensity is not significant: $x^2 = 1.722$, $p = 0.423$.

The results indicate that missingness of smoke is missing at random in relation with sex.

Table 13: Difference in means of missingness of Smoke

variables	M obs	M true	t	p
rest	70.08	68.72	0.779	0.438
age	38.03	40.57	-1.271	0.208
height	174.41	175.12	-0.347	0.731
bmi	24.00	24.85	-1.338	0.188
weight	73.74	76.13	-0.785	0.444

The seven-bar plots in Figures 10 and 11 show how the missing data of smoke is divided into the other columns. The first plot shows us that the female category in sex has fewer missing values in the active column than the male category. The second plot shows almost no difference between missing values in the intensity column for the high and low category, but fewer missing values in the moderate-intensity category. The third plot shows that the missingness of smoke on rest is equally divided with two spikes where rest is lower than 55 and higher than 90. There are no more missing values after these spikes, except for one more spike where the rest is 40. The fourth plot shows that the missingness of smoke on age is equally divided with a spike where age is higher than 60. The fifth plot shows how smaller or higher (than a height of 170-175) the height gets, the more missing values there are in the smoke column, except for when the height is around 205. Then there are close to no missing values. The sixth plot shows that the missingness of smoke on bmi is equally divided except for when bmi is at its lowest or highest. Then there are almost no missing values. The seventh plot shows that the missingness of smoke on weight is equally divided, with one spike in the middle between weight of 70 and 80. There are almost no missing values when weight is at its lowest or highest.
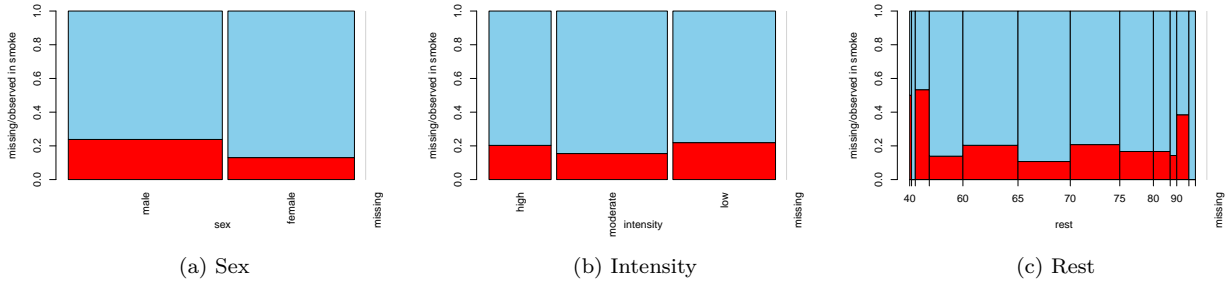

(a) Sex   (b) Intensity   (c) Rest

Figure 10: Looking whether the missingness of smoking is MAR

(a) Age

(b) Height

(c) Bmi

(d) Weight

Figure 11: Looking whether the missingness of smoking is MAR

5 Imputations

After observing the data for missing values, we now try to solve the missingness problem by doing multiple imputations with the package `mice`. Considering our missing data are MAR, mice should work just fine. In order to answer the research question with the imputed data, we will follow the main steps in multiple imputations following van Buuren, 2018, shown in Figure 12. In the first instance, we will use the default settings to impute the missingness, and this will be further elaborated in Section 5.1 section. After the default imputations, we will evaluate the quality of the imputations by examining multiple plots about the convergence and the distribution. The evaluation will be done in Section 5.1.1. After evaluating the default imputations, in Section 5.1.2 we will look at the regressions of the imputed data and compare the outcomes with the regressions of the observed and true data. In Section 5.1.3 we will evaluate the fraction of missing data in the pooled regression analysis. Lastly we will outline how the imputations can be improved by using more sophisticated imputations. This will be discussed in Section 5.1.4.

After doing the Default Imputations, we will try to improve the imputations by using Passive Imputation, described in Section 5.2. We will evaluate the quality of the imputations by examining multiple plots about the convergence and the distribution. The evaluation will be done in Section 5.2.1. After evaluating the default imputations, in Section 5.2.2 we will compare the regression outcomes of the observed, true, and both imputed data. In Section 5.2.3 we will evaluate the fraction of missing data in the pooled regression analysis. Lastly we will outline how the imputations can be improved by using more sophisticated imputations. This will be discussed in Section 5.2.4.

5.1 Default Imputations

As mentioned before, we will first use the default settings of `mice` to impute the missingness. In this section, we will describe them. By default, `mice` will produce 5 imputations and five iterations. An imputation is a replacement value for a missing value. The number of imputations represents the number of imputed data sets that mice creates. Thus, with the default 5 imputations, five datasets are created with imputed data

Incomplete data   Imputed data   Analysis results   Pooled result

Figure 12: Scheme of main steps in multiple imputation

that differ in what values are imputed due to random variation. The random variation also has to do with the uncertainty about what value to impute van Buuren, 2018. In comparison to a single imputation method like mean imputation, the multiple imputation method (obviously) improves the imputation variability. Using a greater amount than just one imputation will result in more accurate standard errors, unbiased estimates, and better confidence intervals (van Buuren, 2018).

The default methods that `mice` uses to impute the missing values are: logreg for smoke, pmm for active, pmm for height, pmm for weight, and pmm for bmi. The method for imputations depends on the measurement level of the target variable. Obviously, there are no methods for *age, sex, intensity, and rest* since they have no missing values. The logreg for smoke uses the Bayesian logistic regression method to impute missing data. The function will calculate the probability of 'Smoker' or 'Non-smoker' for each missing value to impute based on the information from the observations. When comparing the probability to some threshold (probably 50%), the function will impute 'Smoker' for any probability greater than 50% and 'Non-smoker' otherwise. The pmm for active, height, weight, and bmi uses the predictive mean matching method to impute missing data. This "pmm" method starts by defining a linear regression model. The missing values of these variables will be predicted based on observed values (donors) closest to the missing value in the same column. Then, one of these donors is selected randomly, and its value is used to replace the missing value.

The number of iterations, which is 5 by default, represent the amount `mice` will iterate over the variables in each imputation. After imputation, convergence should be met for each variable. If not, one could for example try to enlarge the amount of iterations.

Table 14 provides an overview of the predictor matrix for the imputations. The 1's represent the column variables that will be used as a predictor to impute the row variable. Logically, the 0's will not be used as predictors for imputations. For example, for *bmi* all variables except bmi are used as predictors to impute bmi.

Table 14: Predictor Matrix of Default Imputations

	age	smoke	sex	intensity	active	rest	height	weight	bmi
age	0	1	1	1	1	1	1	1	1
smoke	1	0	1	1	1	1	1	1	1
sex	1	1	0	1	1	1	1	1	1
intensity	1	1	1	0	1	1	1	1	1
active	1	1	1	1	0	1	1	1	1
rest	1	1	1	1	1	0	1	1	1
height	1	1	1	1	1	1	0	1	1
weight	1	1	1	1	1	1	1	0	1
bmi	1	1	1	1	1	1	1	1	0

5.1.1 Evaluating Default Imputations

As can be seen in the trace plots in 13, the imputations, represented as lines, show little to no mix for the variables height, weight, and bmi. This weak convergence suggests that there is not enough information for a solution. The imputed values are thus not plausible. The imputations for smoke and active show a greater convergence. These variablies do thus not require a different imputation method.


Figure 13: Trace plots of the imputations with default options

The density plots in Figure 14 show a variability of the imputed values, in line with the observed values. Futhermore, there appear no impossible values in the densityplots.
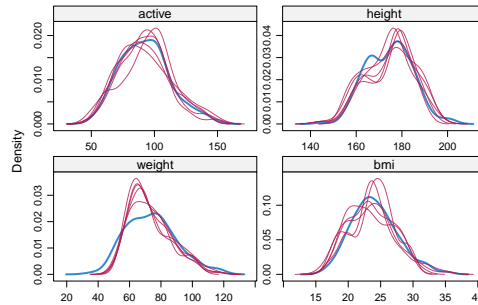

Figure 14: Density plots of the imputed values with default options

16

The stripplot and boxplot for active, as shown in Figure 15 and Figure 16, further increases the validity of the imputation, as the imputed values show a wide variability. For height, and weight, the strip plots show almost no imputations on the lower and higher end. For bmi, the plots show almost no imputations on the higher end. This, however, is in line with the observed data. Based on the strip plots alone, the imputations seem sensible.
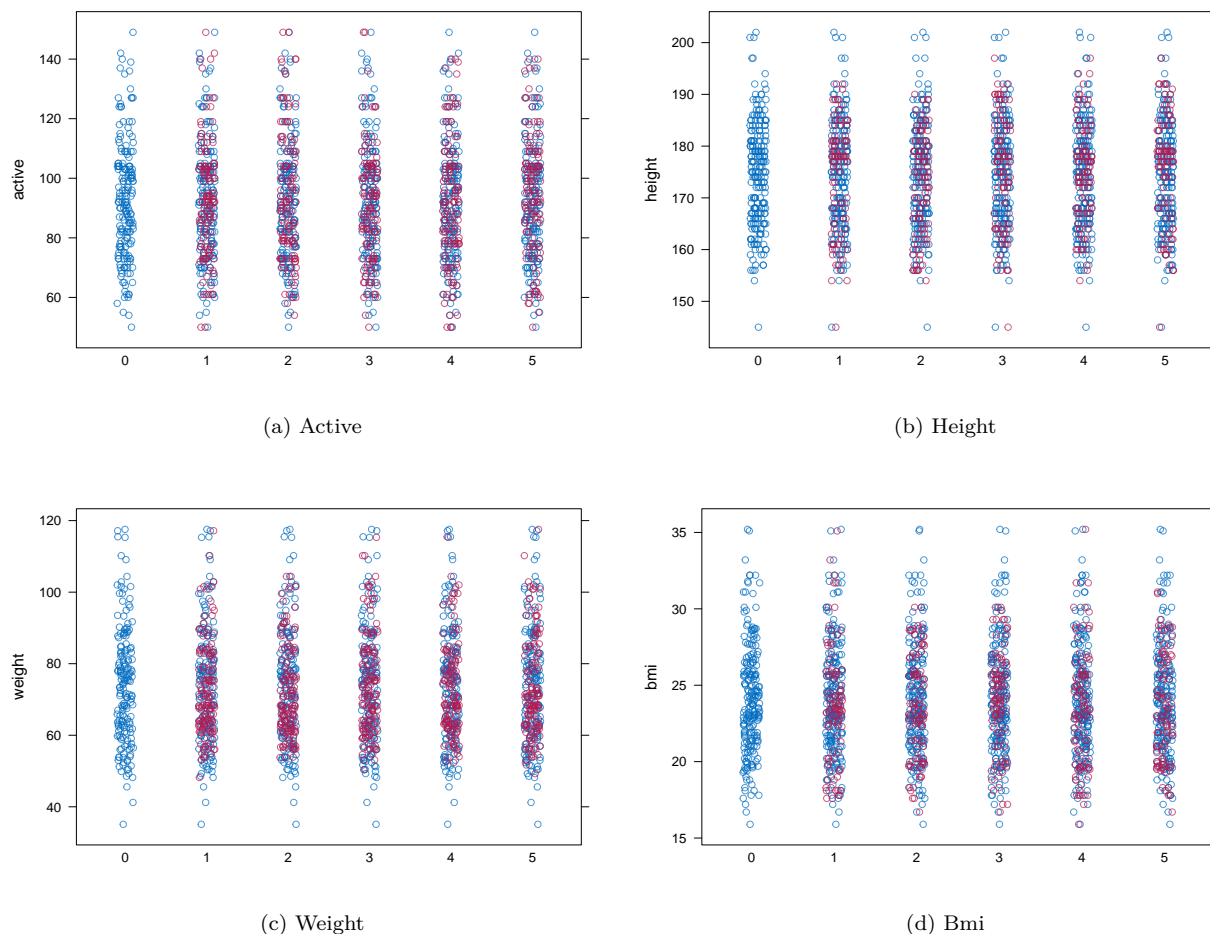

(a) Active

(b) Height

(c) Weight

(d) Bmi

Figure 15: Strip plots of the imputed values with default options
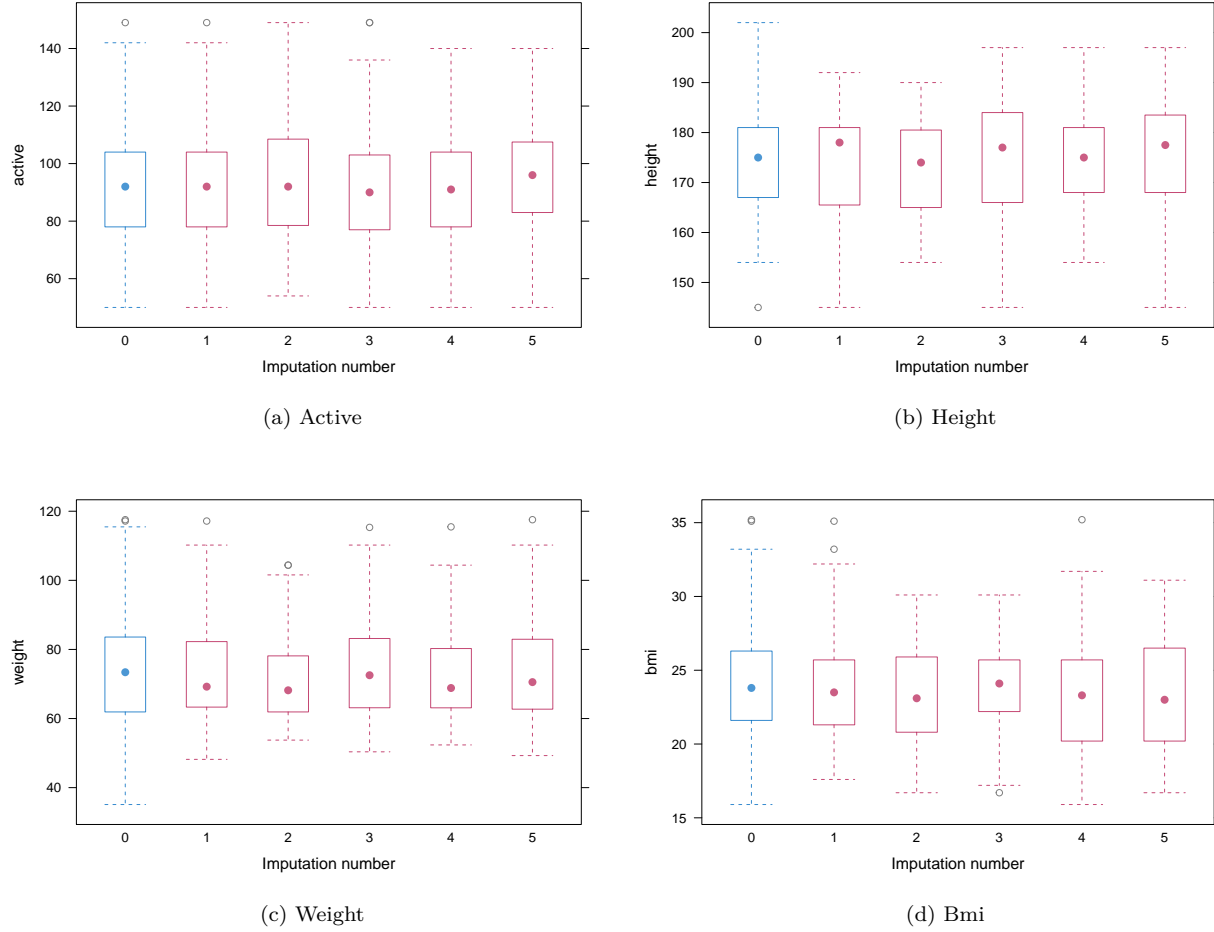
(a) Active

(b) Height

(c) Weight

(d) Bmi

Figure 16: Boxplots of the imputed values with default options

Although the density plots and strip plots for the imputations look acceptable, way me not use the imputed values as the imputation model has shown no convergence. We can hence not make inferences based on the imputed models and need another, more sophisticated model to acquire plausible imputations.

5.1.2 Regressions with Imputed data

In table 15 we present the regression analysis of the observed, imputed and true data. To start with, there are several differences between the observed data and the imputed data. For all predictors in our regression model, the beta coefficients are lower for the imputed data except for *smokeyes*. Thus, the degree of change for *active* has decreased for every predictor, except *smokeyes*. Furthermore, the standard errors around the beta coefficients have decreased for all predictors including the intercept. This suggests that the imputations resulted in more precise estimates for the beta coefficients. Also for every predictor except *smokeyes*, has the p-value increased for the imputed data. An increase in p-value corresponds to lesser significance of a predictor.

Secondly, when comparing the imputed data to the true data, making inferences based on the imputed data is even worse than for the observed data. However, just for the predictor *smokeyes* it has gotten better. One would make the same inferences for *smokeyes* in the true data as for the imputed data. Only the standard error and the p-value for *smokeyes* are still a bit larger in the imputed data than for the true data.

Table 15: Regression analysis of the Observed, Imputed and True data

	Observed Data (N=155)			Imputed Data (N=306)			True Data (N=306)		
	b	SE	p	b	SE	p	b	SE	p
(Intercept)	78.444	14.34	0.000	96.394	13.67	0.000	80.384	9.03	0.000
age	-0.809	0.11	0.000	-0.797	0.08	0.000	-0.883	0.07	0.000
bmi	1.681	0.55	0.003	1.002	0.52	0.077	1.776	0.35	0.000
sexfemale	32.756	20.78	0.117	14.732	16.35	0.371	43.460	14.16	0.002
smokeyes	1.615	2.91	0.580	3.190	2.46	0.202	3.516	1.99	0.078
bmi:sexfemale	-1.131	0.88	0.199	-0.464	0.70	0.508	-1.674	0.60	0.006

5.1.3 Evaluating the fraction of missing data

In Table 16 we present information about the measures of missing values in the imputed data. *RIV* accords for the relative increase in variance, *lambda* is the proportion of the variation in the parameter caused by missing data and *FMI* is basically the same as lambda, but FMI also adjusts for the number of imputed datasets (Heymans & Eekhout, 2019).

In Table 18 we observe that both *bmi* and *active* stands out from the rest. Furthermore, the missing data for *bmi* and *active* accounts for more than 50% of the variation in the parameter. This might has to do with bmi being a derived variable. Later in this document we will try to fix this problem.

Table 16: Missing information in Imputed Data

	riv	lambda	fmi
(Intercept)	1.137	0.532	0.591
age	0.093	0.085	0.095
bmi	1.130	0.531	0.590
sexfemale	0.298	0.229	0.255
smokeyes	0.393	0.282	0.315
bmi:sexfemale	0.332	0.249	0.277

5.1.4 Improving the Imputations

The non-convergence of the imputations for the variables weight, height, and bmi is a structural problem, as bmi is a product of weight and height. This deterministic function is not accounted for in the current imputation model. To account for this function, we will create a new model based on a 'passive imputation' approach. Before we start the next imputations, we specify the relationship between weight, height, and bmi. Then, we alter the predictor matrix, so that the transformed variable is not used as a predictor. As a result, bmi will be predicted based on a function of the already completed variables weight and height.

5.2 Passive Imputations

For the improved imputations we use the technique passive imputation, as described above. After further examination, we decided to predict weight based on a function of bmi and height, as there are more missing values for weight than bmi. To account for this relation, we specify the relation $meth_{weight} = bmi * (\frac{height}{100})^{2}$, and apply this to our imputations. Then, we specify the predictors for height, weight, and bmi. With the setting '0', we ensure that the transformed variables are not being used as predictors for their non-transformed forms. This predictor setting is specified in a variable and will also be applied to the imputation model. The iterations and imputations will remain the same as the previous imputation model.

5.2.1   Evaluating Passive Imputations

As can be seen in the trace plots in Figure 17, the imputations, represented as lines, now intermingle for all predictors. This was not the case for the default imputations as shown in figure 13. Now the lines are also horizontal which is what we like to see. The improved convergence suggests that there is more information for a solution.
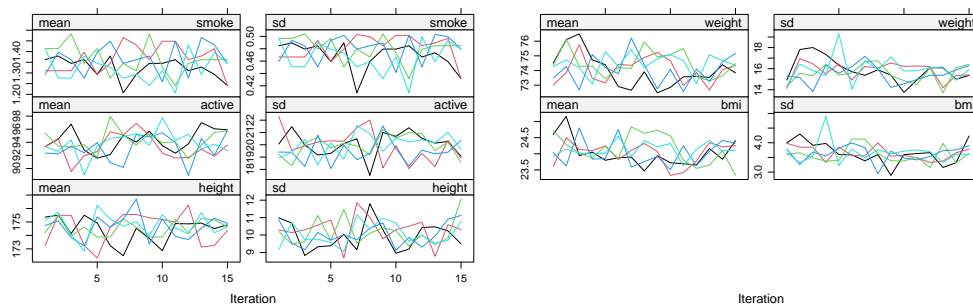

Figure 17: Trace plots of the imputations with Passive Imputation

Also the density plots in Figure 18 show nice outputs as excpected after the trace plots. As mentioned before, the density plots visualize the variability of the imputed values. The imputed values are again plausible.


Figure 18: Density plots of the imputations with Passive Imputation

The strip plots and boxplots for active, are shown in Figure 19 and Figure 20. Both figures look much like the default imputation. However this time, these plots now add value since the trace plots are acceptable. The imputations in the strip plots and boxplots show just a little bit more variability for example in bmi than for the default imputations. The imputations are again in line with the observed data. We consider the passive imputations as sensible.
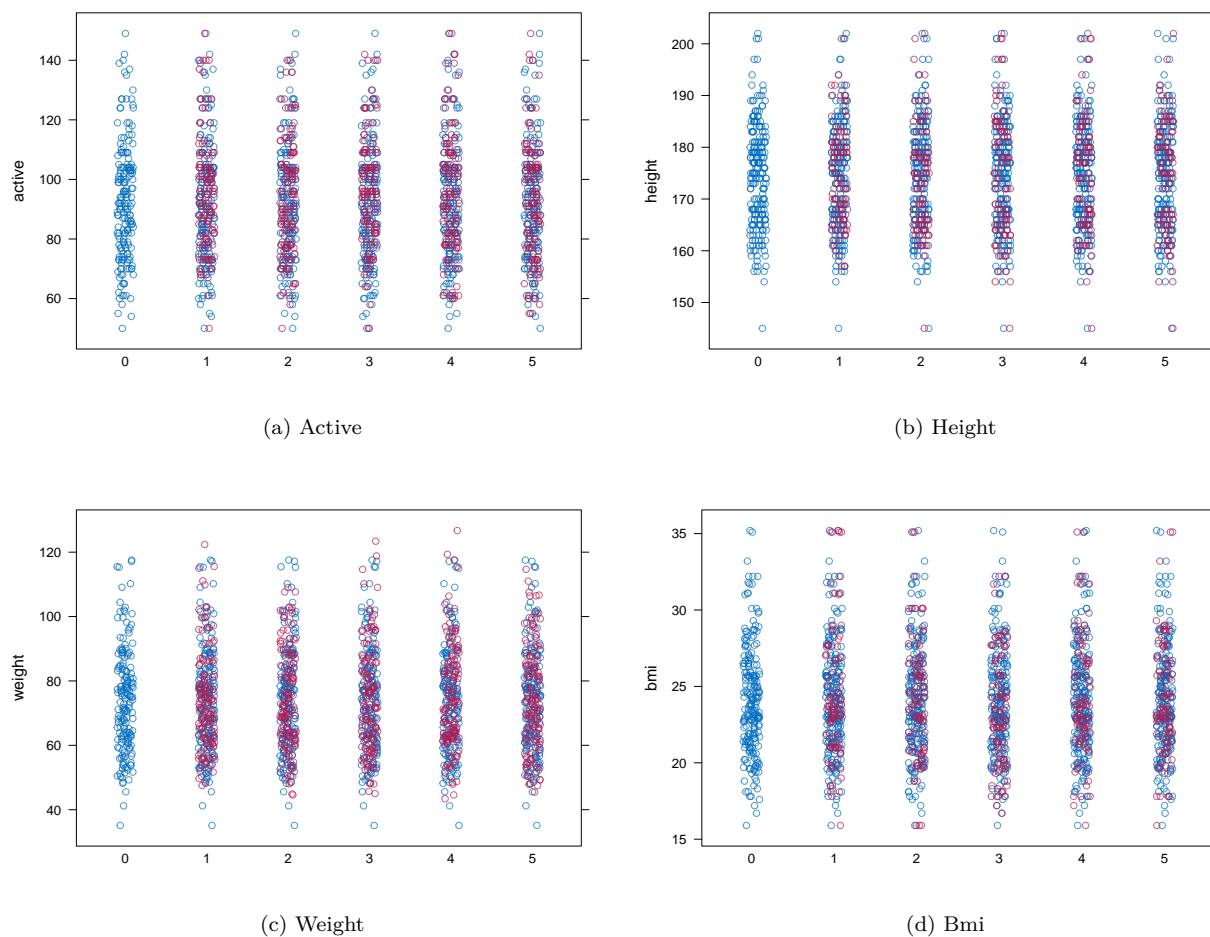

(a) Active

(b) Height

(c) Weight

(d) Bmi

Figure 19: Strip plots of the imputed values with Passive Imputation
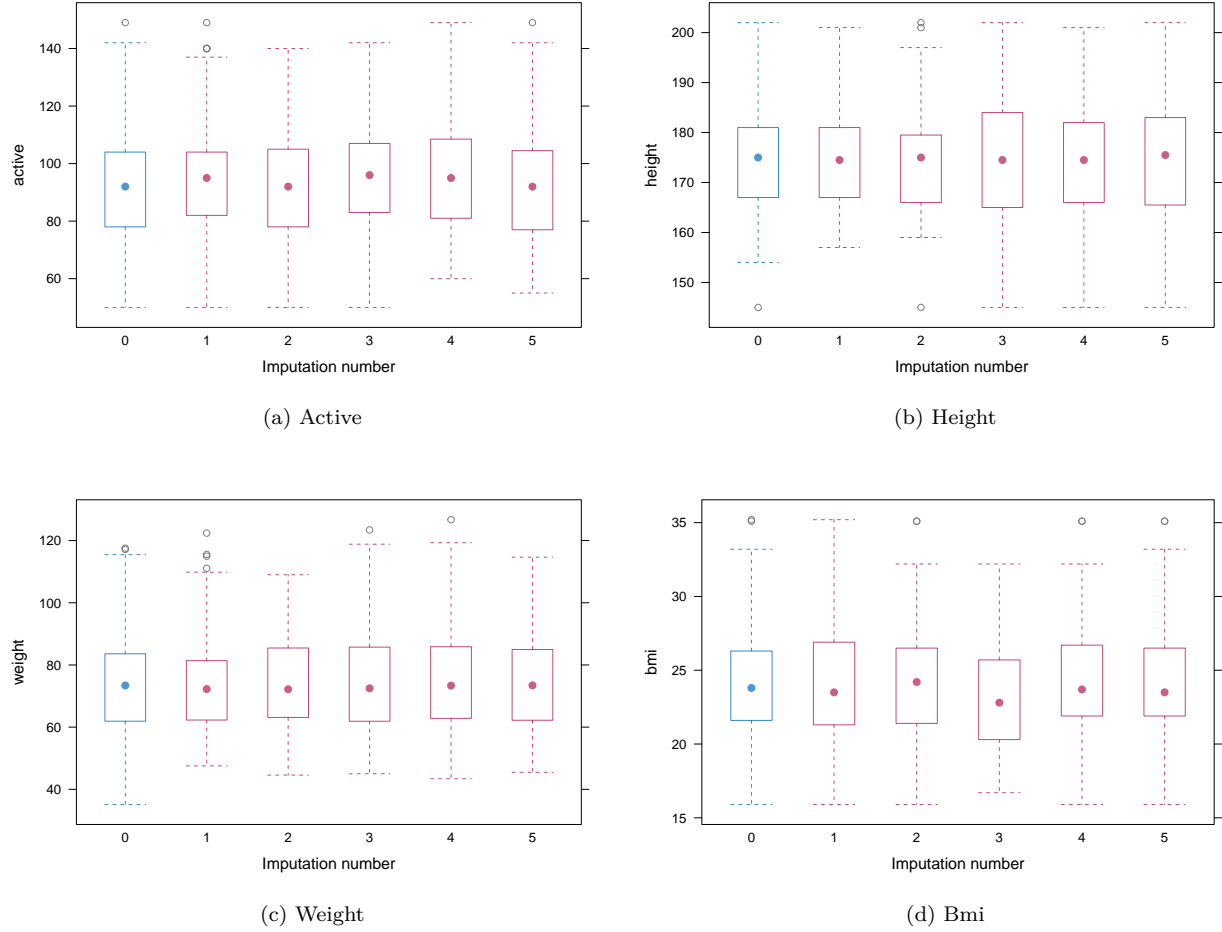
(a) Active

(b) Height

(c) Weight

(d) Bmi

Figure 20: Boxplots of the imputed values with Passive Imputation

5.2.2 Regressions with Passive Imputations

In Table 17 present an extended regression analysis of 15, described in Section 5.1.2, by adding the results of the passive imputations. When we compare the results of the `Default Imputations` with the `Passive Imputations`, the coefficients, the standard errors and the p values for all predictors, but *smokeyes* are more precise in the `Passive Imputations` compared to the `Default Imputations`. However, compared to the `Observed Data`, only the coefficients, standard error and p value of *bmi* seems to have improved by using passive imputation.

Table 17: Regression analysis of the Observed(N=155), Imputed(N=306) and True data(N=306)

	Observed Data			Default Imputation			Passive Imputation			True Data		
	b	SE	p	b	SE	p	b	SE	p	b	SE	p
(Intercept)	78.444	14.34	0.000	96.394	13.67	0.000	90.510	10.22	0.000	80.384	9.03	0.000
age	-0.809	0.11	0.000	-0.797	0.08	0.000	-0.834	0.08	0.000	-0.883	0.07	0.000
bmi	1.681	0.55	0.003	1.002	0.52	0.077	1.334	0.37	0.000	1.776	0.35	0.000
sexfemale	32.756	20.78	0.117	14.732	16.35	0.371	21.940	20.47	0.303	43.460	14.16	0.002
smokeyes	1.615	2.91	0.580	3.190	2.46	0.202	2.253	2.77	0.427	3.516	1.99	0.078
bmi:sexfemale	-1.131	0.88	0.199	-0.464	0.70	0.508	-0.763	0.87	0.396	-1.674	0.60	0.006

5.2.3 Evaluating the Fraction of Missing Data

Table 18 gives us the Fraction of missing information for both the outcomes of the default imputations as the passive imputations. Where we saw high outcomes of *riv*, *lambda*, and *fmi* for *bmi* and *active* with the default imputations, these values drastically decreased after using passive imputation. However, these imputations resulted in substantially higher outcomes of *riv*, *lambda* and *fmi* for *sex*, *smoke* and the interaction of *bmi* and *sex*. These increased values give us the indication that we have to improve the passive imputation.

Table 18: Missing information in Imputed Data

	Default Imputation			Passive Imputation		
	riv	lambda	fmi	riv	lambda	fmi
(Intercept)	1.137	0.532	0.591	0.212	0.175	0.194
age	0.093	0.085	0.095	0.089	0.082	0.091
bmi	1.130	0.531	0.590	0.093	0.085	0.095
sexfemale	0.298	0.229	0.255	1.090	0.522	0.580
smokeyes	0.393	0.282	0.315	0.846	0.458	0.512
bmi:sexfemale	0.332	0.249	0.277	1.124	0.529	0.588

5.2.4 Improving the Imputations

After analyzing the regression model with passive imputations, we conclude that there's quite the room for improvement. For instance, the predictors *smoke*, *sex* and the interaction of *bmi* and *sex* still need a lot of improvement with regard to making accurate inferences out of them. Also, when evaluating the information about the measures of missing values for the passive imputation, we observe values for riv, lambda, and fmi that are to high for the same predictors. Thus, further on this document we will try to finalize these improvements.

5.3 Final Imputations

In this section we will improve the imputations by adding the interaction term which we will analyze to answer our research question in the imputation model, which will be elaborated in Section 5.3.1. In Section 5.3.2 we blablabla. Section 5.3.3. Section 5.3.4. Section 5.3.5. Section 5.3.6

5.3.1 Interaction Term

The first thing we tried to improve the imputations, was adding the interaction term. van Buuren, 2018 contends that not including the complete-data model variables will bias the results towards zero. Thus,

interactions of scientific interest should be included in the imputation model. Therefore we included the interaction of interest (i.e. the interaction of `bmi` and `sex`) in the imputation model.

5.3.2 Methods

Regarding the methods, we use the same methods as described in Section 5.1 (i.e. `pmm` for the continuous variables and `logreg` for the binary variable), and for weight we use the passive imputation ($weight = I(\frac{weight}{(height/100)^2})$), as described in Section 5.2. In our new model, however, we have included the interaction variable. Because the interaction variable consists out of both `bmi` and `sex` we decided to passive impute the missings of this interaction variable as well. We use the method $bmi\_sex = I(bmi * sex)$

5.3.3 Predictor Matrix

Next, we slightly adapted the predictor matrix. Where in the passive imputations, we only excluded `weight` as predictor for `bmi` and `height`, for the final imputations we will make some more adaptations. For starters we did use the `quickpred()` function to include variables who exceed the minimum correlation of 0.1. Predictors who failed to exceed the minimum correlation are eliminated. To prevent circularity we made sure that `weight` is not a predictor for `bmi` and `height`. We did the same for the interaction variable. Table 19 provides the entire predictor matrix we will use for the imputations.

Table 19: Predictor Matrix of Final Imputations

	age	smoke	sex	intensity	active	rest	height	weight	bmi	bmiSex
age	0	0	0	0	0	0	0	0	0	0
smoke	0	0	1	1	1	1	1	1	1	1
sex	0	0	0	0	0	0	0	0	0	0
intensity	0	0	0	0	0	0	0	0	0	0
active	1	1	1	1	0	1	0	1	1	1
rest	0	0	0	0	0	0	0	0	0	0
height	1	1	1	1	1	1	0	0	1	1
weight	1	1	1	1	1	1	1	0	1	1
bmi	1	1	0	0	0	0	1	0	0	0
bmiSex	1	0	0	0	1	0	1	0	0	0

5.3.4 Iterations and Imputations

Because the fraction of missing information are fairly high in the passive imputations, which indicates a substantial missing data problem. Therefore, we decided to increase the amount of imputations. After using trial and error, doing forty imputations led to a the most precise estimates. Furthermore we decided to increase the number of iterations to twenty to get good convergence.

5.3.5 Evaluating the Final Imputations

As can be seen in the trace plots in Figure 21, the imputations, represented as lines, intermingle for all predictors.

Also the density plots in Figure 22 show nice outputs as expected after the trace plots. As mentioned before, the density plots visualize the variability of the imputed values. The imputed values are again plausible.
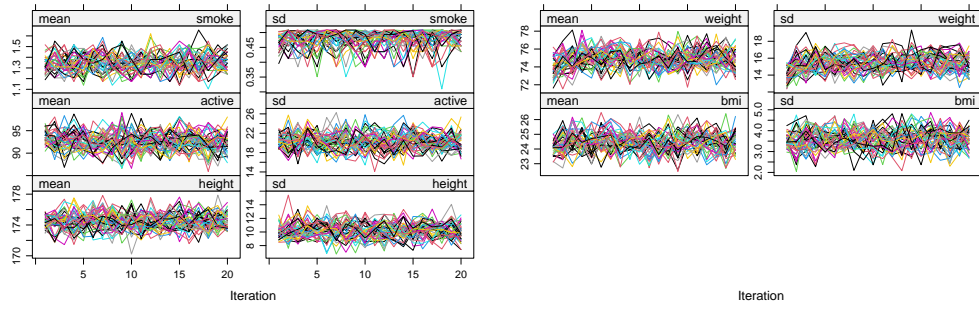
Figure 21: Trace plots of the imputations with Final Imputation


(a) Active

(b) Height

(c) Weight

(d) Bmi

Figure 22: Density plots of the Final Imputations

Table 20: Regression analysis of the Observed(N=155), Imputed and True data(N=306)

	Observed Data			Default Imputation			Passive Imputation			Final Imputation			True Data		
	b	SE	p	b	SE	p	b	SE	p	b	SE	p	b	SE	p
(Intercept)	78.444	14.34	0.000	96.394	13.67	0.000	90.510	10.22	0.000	85.741	12.82	0.000	80.384	9.03	0.000
age	-0.809	0.11	0.000	-0.797	0.08	0.000	-0.834	0.08	0.000	-0.801	0.10	0.000	-0.883	0.07	0.000
bmi	1.681	0.55	0.003	1.002	0.52	0.077	1.334	0.37	0.000	1.397	0.49	0.005	1.776	0.35	0.000
sexfemale	32.756	20.78	0.117	14.732	16.35	0.371	21.940	20.47	0.303	26.410	18.89	0.164	43.460	14.16	0.002
smokeyes	1.615	2.91	0.580	3.190	2.46	0.202	2.253	2.77	0.427	3.194	2.72	0.242	3.516	1.99	0.078
bmi:sexfemale	-1.131	0.88	0.199	-0.464	0.70	0.508	-0.763	0.87	0.396	-0.916	0.80	0.251	-1.674	0.60	0.006

5.3.6 Regressions with Final Imputations

5.3.7 Fraction Missing Information of the Final Imputations

Table 21: Missing information in Imputed Data

	Default Imputation			Passive Imputation			Final Imputation		
	riv	lambda	fmi	riv	lambda	fmi	riv	lambda	fmi
(Intercept)	1.137	0.532	0.591	0.212	0.175	0.194	0.313	0.238	0.249
age	0.093	0.085	0.095	0.089	0.082	0.091	0.335	0.251	0.262
bmi	1.130	0.531	0.590	0.093	0.085	0.095	0.320	0.242	0.253
sexfemale	0.298	0.229	0.255	1.090	0.522	0.580	0.259	0.206	0.216
smokeyes	0.393	0.282	0.315	0.846	0.458	0.512	0.351	0.260	0.271
bmi:sexfemale	0.332	0.249	0.277	1.124	0.529	0.588	0.268	0.211	0.222

6 Conclusion